

Introduction to K-Means Clustering

Machine learning algorithms can be broadly classified into two categories - supervised and unsupervised learning. There are other categories also like semi-supervised learning and reinforcement learning. But, most of the algorithms are classified as supervised or unsupervised learning. The difference between them happens because of presence of target variable. In unsupervised learning, there is no target variable. The dataset only has input variables which describe the data. This is called unsupervised learning.

K-Means clustering is the most popular unsupervised learning algorithm. It is used when we have unlabelled data which is data without defined categories or groups. The algorithm follows an easy or simple way to classify a given data set through a certain number of clusters, fixed apriori. K-Means algorithm works iteratively to assign each data point to one of K groups based on the features that are provided. Data points are clustered based on feature similarity.

K-Means Clustering intuition

K-Means clustering is used to find intrinsic groups within the unlabelled dataset and draw inferences from them. It is based on centroid-based clustering.

Centroid - A centroid is a data point at the centre of a cluster. In centroid-based clustering, clusters are represented by a centroid. It is an iterative algorithm in which the notion of similarity is derived by how close a data point is to the centroid of the cluster. K-Means clustering works as follows:- The K-Means clustering algorithm uses an iterative procedure to deliver a final result. The algorithm requires number of clusters K and the data set as input. The data set is a collection of features for each data point. The algorithm starts with initial estimates for the K centroids. The algorithm then iterates between two steps:-

1. Data assignment step

Each centroid defines one of the clusters. In this step, each data point is assigned to its nearest centroid, which is based on the squared Euclidean distance. So, if c_i is the collection of centroids in set C, then each data point is assigned to a cluster based on minimum Euclidean distance.

1. Centroid update step

In this step, the centroids are recomputed and updated. This is done by taking the mean of all data points assigned to that centroid's cluster.

The algorithm then iterates between step 1 and step 2 until a stopping criteria is met. Stopping criteria means no data points change the clusters, the sum of the distances is minimized or some maximum number of iterations is reached. This algorithm is guaranteed to converge to a result. The result may be a local optimum meaning that

assessing more than one run of the algorithm with randomized starting centroids may give a better outcome.

Choosing the value of K

The K-Means algorithm depends upon finding the number of clusters and data labels for a pre-defined value of K. To find the number of clusters in the data, we need to run the K-Means clustering algorithm for different values of K and compare the results. So, the performance of K-Means algorithm depends upon the value of K. We should choose the optimal value of K that gives us best performance. There are different techniques available to find the optimal value of K. The most common technique is the elbow method which is described below.

The elbow method

The elbow method is used to determine the optimal number of clusters in K-means clustering. The elbow method plots the value of the cost function produced by different values of K.

If K increases, average distortion will decrease. Then each cluster will have fewer constituent instances, and the instances will be closer to their respective centroids. However, the improvements in average distortion will decline as K increases. The value of K at which improvement in distortion declines the most is called the elbow, at which we should stop dividing the data into further clusters.

The problem statement

In this project, We implement K-Means clustering with Python and Scikit-Learn. As mentioned earlier, K-Means clustering is used to find intrinsic groups within the unlabelled dataset and draw inferences from them. I have used Facebook Live Sellers in Thailand Dataset for this project. I implement K-Means clustering to find intrinsic groups within this dataset that display the same status_type behaviour. The status_type behaviour variable consists of posts of a different nature (video, photos, statuses and links).

Dataset description

In this project, I have used Facebook Live Sellers in Thailand Dataset, downloaded from the UCI Machine Learning repository. The dataset can be found at the following url-

<https://archive.ics.uci.edu/ml/datasets/Facebook+Live+Sellers+in+Thailand>

The dataset consists of Facebook pages of 10 Thai fashion and cosmetics retail sellers. The status_type behaviour variable consists of posts of a different nature (video, photos, statuses and links). It also contains engagement metrics of comments, shares and reactions.

Import libraries

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

Ignore warnings

```
import warnings
```

```
warnings.filterwarnings('ignore')
```

Import dataset

```
dataset = pd.read_csv('Live_20210128.csv')
dataset.head()
```

	status_id	status_type	status_published	num_reactions	num_comments
0	1	video	4/22/2018 6:00	529	512
1	2	photo	4/21/2018 22:45	150	0
2	3	video	4/21/2018 6:17	227	236
3	4	photo	4/21/2018 2:29	111	0
4	5	photo	4/18/2018 3:22	213	0

	num_shares	num_likes	num_loves	num_wows	num_hahas	num_sads	\
0	262	432	92	3	1	1	
1	0	150	0	0	0	0	
2	57	204	21	1	1	0	
3	0	111	0	0	0	0	
4	0	204	9	0	0	0	

	num_angrys	Column1	Column2	Column3	Column4
0	0	NaN	NaN	NaN	NaN
1	0	NaN	NaN	NaN	NaN
2	0	NaN	NaN	NaN	NaN
3	0	NaN	NaN	NaN	NaN
4	0	NaN	NaN	NaN	NaN

Exploring Data

```
dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 7050 entries, 0 to 7049
```

```
Data columns (total 16 columns):
```

#	Column	Non-Null Count	Dtype
0	status_id	7050 non-null	int64
1	status_type	7050 non-null	object
2	status_published	7050 non-null	object
3	num_reactions	7050 non-null	int64
4	num_comments	7050 non-null	int64
5	num_shares	7050 non-null	int64
6	num_likes	7050 non-null	int64
7	num_loves	7050 non-null	int64
8	num_wows	7050 non-null	int64
9	num_hahas	7050 non-null	int64
10	num_sads	7050 non-null	int64
11	num_angrys	7050 non-null	int64
12	Column1	0 non-null	float64
13	Column2	0 non-null	float64
14	Column3	0 non-null	float64
15	Column4	0 non-null	float64

```
dtypes: float64(4), int64(10), object(2)
```

```
memory usage: 881.4+ KB
```

```
dataset.describe()
```

	status_id	num_reactions	num_comments	num_shares
count	7050.000000	7050.000000	7050.000000	7050.000000
mean	3525.500000	230.117163	224.356028	40.022553
std	2035.304031	462.625309	889.636820	131.599965
min	1.000000	0.000000	0.000000	0.000000
25%	1763.250000	17.000000	0.000000	0.000000
50%	3525.500000	59.500000	4.000000	0.000000
75%	5287.750000	219.000000	23.000000	4.000000
max	7050.000000	4710.000000	20990.000000	3424.000000

	num_loves	num_wows	num_hahas	num_sads	num_angrys
count	7050.000000	7050.000000	7050.000000	7050.000000	7050.000000
mean	12.728652	1.289362	0.696454	0.243688	0.113191

std	39.972930	8.719650	3.957183	1.597156	0.726812
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000	0.000000	0.000000
75%	3.000000	0.000000	0.000000	0.000000	0.000000
max	657.000000	278.000000	157.000000	51.000000	31.000000

	Column1	Column2	Column3	Column4
count	0.0	0.0	0.0	0.0
mean	NaN	NaN	NaN	NaN
std	NaN	NaN	NaN	NaN
min	NaN	NaN	NaN	NaN
25%	NaN	NaN	NaN	NaN
50%	NaN	NaN	NaN	NaN
75%	NaN	NaN	NaN	NaN
max	NaN	NaN	NaN	NaN

dataset.shape

(7050, 16)

Checking For Missing Values

dataset.isnull().sum()

status_id	0
status_type	0
status_published	0
num_reactions	0
num_comments	0
num_shares	0
num_likes	0
num_loves	0
num_wows	0
num_hahas	0
num_sads	0
num_angrys	0
Column1	7050
Column2	7050
Column3	7050
Column4	7050

dtype: int64

Drop Unused Columns

```
dataset.drop(['Column1', 'Column2', 'Column3', 'Column4'], axis=1,
             inplace=True)
```

```
dataset.head()
```

	status_id	status_type	status_published	num_reactions	num_comments
0	1	video	4/22/2018 6:00	529	512
1	2	photo	4/21/2018 22:45	150	0
2	3	video	4/21/2018 6:17	227	236
3	4	photo	4/21/2018 2:29	111	0
4	5	photo	4/18/2018 3:22	213	0

	num_shares	num_likes	num_loves	num_wows	num_hahas	num_sads
0	262	432	92	3	1	1
1	0	150	0	0	0	0
2	57	204	21	1	1	0
3	0	111	0	0	0	0
4	0	204	9	0	0	0

```
dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 7050 entries, 0 to 7049
```

```
Data columns (total 12 columns):
```

#	Column	Non-Null Count	Dtype
0	status_id	7050 non-null	int64
1	status_type	7050 non-null	object
2	status_published	7050 non-null	object
3	num_reactions	7050 non-null	int64
4	num_comments	7050 non-null	int64
5	num_shares	7050 non-null	int64
6	num_likes	7050 non-null	int64
7	num_loves	7050 non-null	int64
8	num_wows	7050 non-null	int64
9	num_hahas	7050 non-null	int64
10	num_sads	7050 non-null	int64

```
11 num_angrys          7050 non-null    int64
dtypes: int64(10), object(2)
memory usage: 661.1+ KB
```

```
dataset.describe()
```

	status_id	num_reactions	num_comments	num_shares
num_likes \				
count	7050.000000	7050.000000	7050.000000	7050.000000
7050.000000				
mean	3525.500000	230.117163	224.356028	40.022553
215.043121				
std	2035.304031	462.625309	889.636820	131.599965
449.472357				
min	1.000000	0.000000	0.000000	0.000000
0.000000				
25%	1763.250000	17.000000	0.000000	0.000000
17.000000				
50%	3525.500000	59.500000	4.000000	0.000000
58.000000				
75%	5287.750000	219.000000	23.000000	4.000000
184.750000				
max	7050.000000	4710.000000	20990.000000	3424.000000
4710.000000				

	num_loves	num_wows	num_hahas	num_sads	num_angrys
count	7050.000000	7050.000000	7050.000000	7050.000000	7050.000000
mean	12.728652	1.289362	0.696454	0.243688	0.113191
std	39.972930	8.719650	3.957183	1.597156	0.726812
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000	0.000000	0.000000
75%	3.000000	0.000000	0.000000	0.000000	0.000000
max	657.000000	278.000000	157.000000	51.000000	31.000000

Explore status_id variable

```
# view the labels in the variable
```

```
dataset['status_id'].unique()
```

```
array([ 1,  2,  3, ..., 7048, 7049, 7050], dtype=int64)
```

```
# view how many different types of variables are there  
len(dataset['status_id'].unique())
```

7050

We can see that there are 6997 unique labels in the status_id variable. The total number of instances in the dataset is 7050. So, it is approximately a unique identifier for each of the instances. Thus this is not a variable that we can use. Hence, I will drop it.

Explore status_published variable

```
# view the labels in the variable
```

```
dataset['status_published'].unique()  
  
array(['4/22/2018 6:00', '4/21/2018 22:45', '4/21/2018 6:17', ...,  
      '9/21/2016 23:03', '9/20/2016 0:43', '9/10/2016 10:30'],  
      dtype=object)
```

```
# view how many different types of variables are there
```

```
len(dataset['status_published'].unique())
```

6913

Again, we can see that there are 6913 unique labels in the status_published variable. The total number of instances in the dataset is 7050. So, it is also approximately a unique identifier for each of the instances. Thus this is not a variable that we can use. Hence, I will drop it also.

Explore status_type variable

```
# view the labels in the variable
```

```
dataset['status_type'].unique()  
  
array(['video', 'photo', 'link', 'status'], dtype=object)
```

```
# view how many different types of variables are there
```

```
len(dataset['status_type'].unique())
```

4

We can see that there are 4 categories of labels in the status_type variable.

Drop status_id and status_published variable from the dataset

```
dataset.drop(['status_id', 'status_published'], axis=1, inplace=True)
```


View the summary of dataset again

```
dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7050 entries, 0 to 7049
Data columns (total 10 columns):
 #   Column              Non-Null Count  Dtype  
---  -
 0   status_type         7050 non-null   object 
 1   num_reactions        7050 non-null   int64  
 2   num_comments         7050 non-null   int64  
 3   num_shares           7050 non-null   int64  
 4   num_likes            7050 non-null   int64  
 5   num_loves            7050 non-null   int64  
 6   num_wows             7050 non-null   int64  
 7   num_hahas            7050 non-null   int64  
 8   num_sads             7050 non-null   int64  
 9   num_angrys          7050 non-null   int64  
dtypes: int64(9), object(1)
memory usage: 550.9+ KB
```

```
dataset.describe()
```

	num_reactions	num_comments	num_shares	num_likes
count	7050.000000	7050.000000	7050.000000	7050.000000
mean	230.117163	224.356028	40.022553	215.043121
std	462.625309	889.636820	131.599965	449.472357
min	0.000000	0.000000	0.000000	0.000000
25%	17.000000	0.000000	0.000000	17.000000
50%	59.500000	4.000000	0.000000	58.000000
75%	219.000000	23.000000	4.000000	184.750000
max	4710.000000	20990.000000	3424.000000	4710.000000

	num_wows	num_hahas	num_sads	num_angrys
count	7050.000000	7050.000000	7050.000000	7050.000000
mean	1.289362	0.696454	0.243688	0.113191
std	8.719650	3.957183	1.597156	0.726812
min	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000	0.000000

75%	0.000000	0.000000	0.000000	0.000000
max	278.000000	157.000000	51.000000	31.000000

Declare feature vector and target variable

```
X = dataset
```

```
y = dataset['status_type']
```

Convert categorical variable into integers

```
from sklearn.preprocessing import LabelEncoder
```

```
le = LabelEncoder()
```

```
X['status_type'] = le.fit_transform(X['status_type'])
```

```
y = le.transform(y)
```

```
X.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7050 entries, 0 to 7049
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   status_type           7050 non-null   int32
1   num_reactions         7050 non-null   int64
2   num_comments          7050 non-null   int64
3   num_shares            7050 non-null   int64
4   num_likes             7050 non-null   int64
5   num_loves             7050 non-null   int64
6   num_wows              7050 non-null   int64
7   num_hahas            7050 non-null   int64
8   num_sads              7050 non-null   int64
9   num_angrys            7050 non-null   int64
dtypes: int32(1), int64(9)
memory usage: 523.4 KB
```

```
X.head()
```

	status_type	num_reactions	num_comments	num_shares	num_likes
0	3	529	512	262	432
92	1	150	0	0	150
2	3	227	236	57	204
21	1	111	0	0	111

```
0
4          1          213          0          0          204
9
```

```
    num_wows  num_hahas  num_sads  num_angrys
0          3          1          1          0
1          0          0          0          0
2          1          1          0          0
3          0          0          0          0
4          0          0          0          0
```

Feature Scaling

```
cols = X.columns
```

```
from sklearn.preprocessing import MinMaxScaler
```

```
ms = MinMaxScaler()
```

```
X = ms.fit_transform(X)
```

```
X = pd.DataFrame(X, columns=[cols])
```

```
X.head()
```

```
    status_type  num_reactions  num_comments  num_shares  num_likes
num_loves \
0    1.000000    0.112314    0.024393    0.076519    0.091720
0.140030
1    0.333333    0.031847    0.000000    0.000000    0.031847
0.000000
2    1.000000    0.048195    0.011243    0.016647    0.043312
0.031963
3    0.333333    0.023567    0.000000    0.000000    0.023567
0.000000
4    0.333333    0.045223    0.000000    0.000000    0.043312
0.013699
```

```
    num_wows  num_hahas  num_sads  num_angrys
0  0.010791  0.006369  0.019608    0.0
1  0.000000  0.000000  0.000000    0.0
2  0.003597  0.006369  0.000000    0.0
3  0.000000  0.000000  0.000000    0.0
4  0.000000  0.000000  0.000000    0.0
```

K-Means model with two clusters

```
from sklearn.cluster import KMeans
```

```
kmeans = KMeans(n_clusters=2, random_state=0)
```

```
kmeans.fit(X)
KMeans(n_clusters=2, random_state=0)
```

K-Means model parameters study

```
kmeans.cluster_centers_
array([[3.28506857e-01, 3.90710874e-02, 7.54854864e-04, 7.53667113e-
04,
        3.85438884e-02, 2.17448568e-03, 2.43721364e-03, 1.20039760e-
03,
        2.75348016e-03, 1.45313276e-03],
       [9.54921576e-01, 6.46330441e-02, 2.67028654e-02, 2.93171709e-
02,
        5.71231462e-02, 4.71007076e-02, 8.18581889e-03, 9.65207685e-
03,
        8.04219428e-03, 7.19501847e-03]])
```

The KMeans algorithm clusters data by trying to separate samples in n groups of equal variances, minimizing a criterion known as inertia, or within-cluster sum-of-squares. Inertia, or the within-cluster sum of squares criterion, can be recognized as a measure of how internally coherent clusters are. The k-means algorithm divides a set of N samples X into K disjoint clusters C , each described by the mean j of the samples in the cluster. The means are commonly called the cluster centroids. The K-means algorithm aims to choose centroids that minimize the inertia, or within-cluster sum of squared criterion. Inertia is not a normalized metric.

The lower values of inertia are better and zero is optimal.

But in very high-dimensional spaces, euclidean distances tend to become inflated (this is an instance of curse of dimensionality).

Running a dimensionality reduction algorithm such as PCA prior to k-means clustering can alleviate this problem and speed up the computations.

We can calculate model inertia as follows:-

```
kmeans.inertia_
237.7572640441955
```

Check quality of weak classification by the model

```
labels = kmeans.labels_

# check how many of the samples were correctly labeled
correct_labels = sum(y == labels)
```

```
print("Result: %d out of %d samples were correctly labeled." %  
(correct_labels, y.size))
```

Result: 63 out of 7050 samples were correctly labeled.


```
print('Accuracy score: {0:0.2f}'.  
format(correct_labels/float(y.size)))
```

Accuracy score: 0.01

We have achieved a weak classification accuracy of 1% by our unsupervised model.

Use elbow method to find optimal number of clusters

```
from sklearn.cluster import KMeans  
cs = []  
for i in range(1, 11):  
    kmeans = KMeans(n_clusters = i, init = 'k-means++', max_iter =  
300, n_init = 10, random_state = 0)  
    kmeans.fit(X)  
    cs.append(kmeans.inertia_)  
plt.plot(range(1, 11), cs)  
plt.title('The Elbow Method')  
plt.xlabel('Number of clusters')  
plt.ylabel('CS')  
plt.show()
```



By the above plot, we can see that there is a kink at k=2.

Hence $k=2$ can be considered a good number of the cluster to cluster this data.

But, we have seen that I have achieved a weak classification accuracy of 1% with $k=2$.

I will write the required code with $k=2$ again for convinience.

```
from sklearn.cluster import KMeans

kmeans = KMeans(n_clusters=2, random_state=0)

kmeans.fit(X)

labels = kmeans.labels_

# check how many of the samples were correctly labeled

correct_labels = sum(y == labels)

print("Result: %d out of %d samples were correctly labeled." %
      (correct_labels, y.size))

print('Accuracy score: {0:0.2f}'.
      format(correct_labels/float(y.size)))

Result: 63 out of 7050 samples were correctly labeled.
Accuracy score: 0.01
```

So, our weak unsupervised classification model achieved a very weak classification accuracy of 1%.

I will check the model accuracy with different number of clusters.

K-Means model with different clusters

K-Means model with 3 clusters

```
kmeans = KMeans(n_clusters=3, random_state=0)

kmeans.fit(X)

# check how many of the samples were correctly labeled
labels = kmeans.labels_

correct_labels = sum(y == labels)
print("Result: %d out of %d samples were correctly labeled." %
      (correct_labels, y.size))
print('Accuracy score: {0:0.2f}'.
      format(correct_labels/float(y.size)))
```

Result: 6 out of 7050 samples were correctly labeled.
Accuracy score: 0.00

K-Means model with 4 clusters

```
kmeans = KMeans(n_clusters=4, random_state=0)
```

```
kmeans.fit(X)
```

```
# check how many of the samples were correctly labeled
```

```
labels = kmeans.labels_
```

```
correct_labels = sum(y == labels)
```

```
print("Result: %d out of %d samples were correctly labeled." %  
(correct_labels, y.size))
```

```
print('Accuracy score: {0:0.2f}'.  
format(correct_labels/float(y.size)))
```

Result: 51 out of 7050 samples were correctly labeled.
Accuracy score: 0.01

We have achieved a relatively high accuracy of 62% with k=4.

1. Results and conclusion In this project, I have implemented the most popular unsupervised clustering technique called K-Means Clustering.

I have applied the elbow method and find that k=2 (k is number of clusters) can be considered a good number of cluster to cluster this data.

I have find that the model has very high inertia of 237.7572. So, this is not a good model fit to the data.

I have achieved a weak classification accuracy of 1% with k=2 by our unsupervised model.

So, I have changed the value of k and find relatively higher classification accuracy of 62% with k=4.

Hence, we can conclude that k=4 being the optimal number of clusters.